

# Codex Skills Audit Report

Audit date: April 19, 2026  
Repository: /Users/maggielerman/Github/codex-skills

## Summary

| Check                 | Result                     | Notes                                                                                                                    |
|-----------------------|----------------------------|--------------------------------------------------------------------------------------------------------------------------|
| Suite completeness    | Pass                       | All 20 skills in skills/SUITE_SKILLS.txt exist in the repo and in local Codex state.                                     |
| Main repo purpose     | Clear                      | Curated portable skill suite plus selected custom plugin backups.                                                        |
| Biggest strengths     | Strong                     | Docs/governance stack is coherent; execution stack is complementary; plugin backups are separated cleanly.               |
| Biggest problems      | Real                       | Portability drift, lifecycle metadata drift, duplicate tranche-orchestrator ownership, thin tests, generated junk files. |
| Best evaluator result | design-system-refine A/100 | Best overall target from Plugin Eval.                                                                                    |
| Most important action | Hardening                  | Fix portability and de-duplicate tranche orchestration before doing broader cleanup.                                     |

## Suite Skills Table

| Skill                        | Cluster    | Primary job                                                      | Works with                                               | Lifecycle | Plugin Eval | Keep or absorb | Key issues                                        |
|------------------------------|------------|------------------------------------------------------------------|----------------------------------------------------------|-----------|-------------|----------------|---------------------------------------------------|
| ML-catalog-review            | Assessment | Shopify-first catalog review workflow and artifact loop          | ML-docs-system-scaffold optional overlay                 | active    | C/72        | Keep separate  | Large description; no helper tests                |
| ML-cluster-assignment        | Execution  | Split work into safe parallel clusters and bootstrap worktrees   | ML-project-tranche-orchestrator                          | active    | B/86        | Keep separate  | Healthy overall; could add tests for shell helper |
| ML-codex-to-claude-handoff   | Utilities  | Generate Claude repo rules from Codex instructions               | ML-docs-system-scaffold optional                         | active    | C/77        | Keep separate  | Hard-coded paths; tests missing                   |
| ML-collaborative-walkthrough | Governance | Structured sign-off walkthrough for in-review projects           | ML-docs-system-scaffold, ML-project-tranche-orchestrator | active    | B/86        | Keep separate  | Low risk; mostly evaluator naming false positive  |
| ML-continue-until-blocked    | Execution  | Continuous autonomous execution once scope is already approved   | ML-project-tranche-orchestrator                          | active    | D/67        | Keep separate  | Very long prompt; token heavy                     |
| ML-docs-evidence-backfill    | Docs core  | Retrofit fragmented evidence into canonical docs evidence system | ML-docs-system-scaffold                                  | active    | C/72        | Keep separate  | Portability links; trigger text is heavy          |

| Skill                               | Cluster    | Primary job                                                             | Works with                              | Lifecycle                                      | Plugin Eval | Keep or absorb            | Key issues                                                       |
|-------------------------------------|------------|-------------------------------------------------------------------------|-----------------------------------------|------------------------------------------------|-------------|---------------------------|------------------------------------------------------------------|
| ML-docs-product-pack                | Docs core  | Deprecated shim pointing to scaffold flow                               | ML-docs-system-scaffold                 | deprecat ed                                    | C/81        | Absorb into scaffold      | Deprecated shim only                                             |
| ML-docs-project-dashboa rd-legacy   | Docs core  | Standalone retrofit for project dashboard output                        | ML-docs-system-scaffold                 | active in metadata, reads legacy in skill text | D/63        | Absorb into scaffold      | Metadata drift; should not still read as active                  |
| ML-docs-project-lifecyc le-statuses | Docs core  | Retrofit older repos with newer lifecycle states                        | ML-docs-system-scaffold                 | legacy                                         | D/68        | Keep legacy, then archive | Migration-only value                                             |
| ML-docs-system-scaffold             | Docs core  | Create or standardize docs/governance/evidence system                   | evidence-backfill, walkthrough, tranche | active                                         | F/49        | Keep core                 | Portability issues; high token budget; dashboard path dependency |
| ML-env-bootstrap-sync               | Utilities  | Standardize env/bootstrap and worktree sync patterns                    | ML-docs-system-scaffold optional        | active                                         | B/86        | Keep separate             | Low risk; mostly evaluator naming false positive                 |
| ML-github-docs-tracking-sync        | Utilities  | Optional docs-to-GitHub tracking sync                                   | ML-docs-system-scaffold optional        | active                                         | D/63        | Maybe absorb later        | Niche utility; helper complexity                                 |
| ML-hyphenomenon-project-intake      | Assessment | Create import-ready intake dossier plus screenshots                     | -                                       | active                                         | D/68        | Keep separate             | Helper tests missing                                             |
| ML-plan-signoff-review              | Governance | Final cross-check before plan approval                                  | ML-project-gove rnance-audit            | active                                         | B/86        | Keep separate             | Low risk; mostly evaluator naming false positive                 |
| ML-project-governance-audit         | Governance | Audit roadmap order, blockers, and checkpoint hygiene                   | tranche, walkthrough, signoff           | active                                         | C/81        | Keep separate             | Description could be sharper                                     |
| ML-project-tranche-orchestrator     | Execution  | Primary-agent tranche selection, delegation, review, and docs writeback | docs scaffold, continue, cluster        | active                                         | F/54        | Keep, but de-dupe plugin  | Duplicate with plugin; helper tests missing                      |
| ML-repo-implementation-review       | Assessment | Compare 2-3 repos and generate report artifact                          | -                                       | active                                         | F/54        | Keep separate             | Large deferred bundle; helper complexity                         |
| ML-shopify-app-scaffold             | Utilities  | Scaffold Shopify app shell with reusable overlays                       | -                                       | active                                         | D/63        | Keep separate             | Hard-coded paths; helper complexity                              |
| ML-user-journey-audit               | Assessment | Route and UX journey audit with PDF output                              | ML-ux-ui-bug-in take optional           | active                                         | D/63        | Keep separate             | Hard-coded paths; tests missing                                  |
| ML-ux-ui-bug-intake                 | Utilities  | Set up lightweight UX/UI bug intake workflow                            | ML-user-journey-audit optional          | active                                         | C/77        | Keep separate             | openai.yaml schema mismatch; hard-coded path                     |

## Plugin And Plugin-Skill Table

| Item                         | Type         | Primary job                                    | Works with                                  | Plugin Eval | Keep or absorb           | Key issues                                             |
|------------------------------|--------------|------------------------------------------------|---------------------------------------------|-------------|--------------------------|--------------------------------------------------------|
| design-system-first          | Plugin       | Frontend starter and refinement bundle         | design-system-starter, design-system-refine | F/31        | Keep as plugin           | Bundle score low because metadata and budget are heavy |
| design-system-starter        | Plugin skill | Scaffold design-system-first Next.js starters  | design-system-refine                        | C/82        | Keep                     | Good candidate for test hardening                      |
| design-system-refine         | Plugin skill | Refine an existing design-system-first starter | design-system-starter                       | A/100       | Keep                     | Best evaluator result in audit                         |
| project-tranche-orchestrator | Plugin       | Plugin packaging of tranche orchestration      | suite tranche skill                         | C/73        | De-dupe with suite skill | Should not be a second logic source                    |
| project-tranche-orchestrator | Plugin skill | Plugin-contained tranche orchestrator skill    | suite ML-project-tranche-orchestrator       | C/81        | Absorb or sync           | Use one canonical source                               |

## Cross-Cutting Findings

| Finding                             | Impact | Affects                                                                                                                                       | Recommended action                                                                                   |
|-------------------------------------|--------|-----------------------------------------------------------------------------------------------------------------------------------------------|------------------------------------------------------------------------------------------------------|
| Portability drift                   | High   | ML-docs-system-scaffold, ML-codex-to-claude-handoff, ML-user-journey-audit, ML-shopify-app-scaffold, ML-ux-ui-bug-intake, and some docs links | Replace hard-coded .codex and user-specific paths with skill-relative guidance                       |
| Lifecycle metadata drift            | High   | ML-docs-project-dashboard-legacy                                                                                                              | Mark as deprecated or legacy in skills/SUITE_METADATA.json and regenerate catalog                    |
| Duplicate source of truth           | High   | ML-project-tranche-orchestrator plus plugin copy                                                                                              | Choose one canonical source and sync or generate the other                                           |
| Catalog parser mismatch             | Medium | ML-ux-ui-bug-intake                                                                                                                           | Normalize agents/openai.yaml or extend scripts/build_catalog.py to support both metadata shapes      |
| Generated junk in bundles           | Medium | multiple skills/plugins                                                                                                                       | Remove .DS_Store, __pycache__, .pyc; update .gitignore                                               |
| Thin test coverage                  | Medium | most helper-script skills                                                                                                                     | Add tests for tranche helpers, docs scaffold helpers, repo review helpers, and design-system starter |
| Default Plugin Eval rubric mismatch | Medium | almost all ML - * skills                                                                                                                      | Consider a custom metric pack so intentional ML - * naming stops dominating scores                   |

## Local Skills Not Mirrored In This Repo

| Status                                   | Count | Notes                                                                                      |
|------------------------------------------|-------|--------------------------------------------------------------------------------------------|
| Extra local skills not in repo allowlist | 26    | This is not a failure if the repo remains intentionally curated rather than a full mirror. |

### Local-only skills

atlas, cloudflare-deploy, codex-primary-runtime, doc, gh-address-comments, gh-fix-ci, imagegen, jupyter-notebook, netlify-deploy, notion-knowledge-capture, notion-meeting-intelligence, notion-research-documentation, notion-spec-to-implementation, openai-docs, pdf, playwright, playwright-interactive, render-deploy, screenshot, security-best-practices, sora, speech, spreadsheet, transcribe, vercel-deploy, wrapped

## Priority Actions

| Priority | Action                                                | Why first                                                                                |
|----------|-------------------------------------------------------|------------------------------------------------------------------------------------------|
| 1        | Fix portability in flagship skills                    | Directly supports the repo's core promise that skills remain portable standalone bundles |
| 2        | De-dupe tranche orchestrator between suite and plugin | Prevents drift in one of the most important workflow skills                              |
| 3        | Correct lifecycle metadata for dashboard legacy skill | Makes generated catalog and repo status truthful again                                   |
| 4        | Remove generated junk files                           | Quick cleanup that improves portability and professionalism                              |
| 5        | Add tests for highest-value helper scripts            | Reduces regression risk in the most automation-heavy skills                              |
| 6        | Add custom Plugin Eval metric pack                    | Improves future audits by reducing low-signal naming penalties                           |